

Data Quality Assessment for the Formula 1 Reconciled Database

Giuseppe Rudi

May 2, 2026

Contents

Introduction	3
Role of Data Quality Assessment	3
General Data Quality Assessment	3
Theoretical Background	4
Assessment Methodology	5
Score Interpretation	5
Main Output Files	5
Table-by-Table Data Quality Plan	6
Season Table	6
Circuit Table	6
Driver Table	7
Team Table	8
Grand Prix Table	8
Session Table	9
Result Table	9
Lap Table	9
Weather Table	10
Track Status Table	11
Planned Visual Outputs	12
Overall Data Quality Score by Table	12
Data Quality Score by Dimension	13
Heatmap of Table-Dimension Scores	13
Missing Values Summary	14
Referential Integrity Issues	14
Focused Missing Values Analysis	15
Missing Values and Completeness	15
MCAR, MAR, and MNAR Classification	16
Expected NULL Values	16
Problematic NULL Values	17
Focused Outlier Detection	17
IQR Method	17
Modified Z-score	18
Consensus and Multi-method Flag	18
Domain Interpretation	18

Optional LLM Support	19
Not Part of the Deterministic Pipeline	19
Used Only for Explanation	20
Manually Validated Output	20
Conclusion	20

Introduction

This document describes the Data Quality Assessment phase applied to the Formula 1 reconciled database. The goal of this phase is not to clean the data directly, but to evaluate the current quality level of the data stored in PostgreSQL before applying cleaning rules and before loading the final Data Warehouse.


Since the Data Warehouse and the final analytical dashboards depend on the quality of the reconciled level, a systematic assessment is required before continuing with the cleaning and ETL phases.

In this project, Data Quality Assessment is implemented as a deterministic and reproducible Python process that reads the reconciled tables from PostgreSQL, applies a set of quality checks, and produces scorecards and row-level issue files.

Role of Data Quality Assessment

The Data Quality Assessment phase acts as an audit checkpoint between the reconciled database and the cleaned analytical data. Its purpose is to detect data problems before they propagate to the Data Warehouse, the star schema, and the final visual analytics layer.

The project pipeline can be summarized as follows:



General Data Quality Assessment

It is important to distinguish three different activities: Data Quality Assessment, Data Cleaning, and Analytical Filtering. They are related, but they have different goals.

Activity	Goal	Example in this project
Data Quality Assessment	Measure the current quality level and identify problems.	Count missing values in <code>lap</code> , detect duplicate keys, find invalid sector times.
Data Cleaning	Correct, remove, or standardize data quality problems.	Remove deleted or inaccurate laps, standardize textual values.
Analytical Filtering	Select the subset of valid data useful for the final story.	Keep only Ferrari, Red Bull, and Mercedes; keep only dry conditions.

Theoretical Background

Data quality is a multidimensional concept. For this reason, a single global value is not sufficient to understand whether a dataset is reliable.

The assessment in this project follows the ISO 25012 perspective introduced in the course material. The six main dimensions considered are:

- Completeness
- Uniqueness
- Validity
- Consistency
- Accuracy
- Timeliness

Since this project is based on a relational PostgreSQL reconciled database, an additional practical dimension is also considered:

- Referential Integrity

This additional dimension is essential before adding SQL foreign key constraints and before designing the ETL toward the Data Warehouse.

Completeness Completeness measures whether required values are present. In practical terms, a completeness check verifies the percentage of non-null values for attributes that are necessary to identify, connect, or analyze the records.

In this project, completeness is applied differently depending on the role of each attribute. Structural identifiers such as `lap_id`, `session_id`, `driver_id`, and `team_id` are expected to be complete. Analytical attributes such as `lap_time_ms` may be missing in some special cases and must be evaluated more carefully.

Formula used conceptually:

$$Completeness(A) = \frac{\# non-null\ values\ in\ A}{\# total\ rows}$$

Uniqueness Uniqueness measures whether keys and candidate keys are free from duplicates. This dimension is crucial because the reconciled database must later support primary key constraints and reliable joins.

In this project, uniqueness is checked both on surrogate identifiers and on natural keys. For example, `lap_id` must be unique if it is used as surrogate key, while the combination `session_id + driver_id + lap_number` should also identify a lap naturally.

Formula used conceptually:

$$Uniqueness(K) = \frac{\# distinct\ key\ values}{\# total\ key\ values}$$

Validity Validity measures whether values conform to allowed domains, formats, and business rules. A value may be present and unique, but still invalid if it violates the expected domain.

Consistency Consistency measures whether values are coherent with each other. This type of check usually compares two or more attributes in the same table, or values stored in different related tables.

Accuracy and Plausibility Accuracy measures whether values reflect the real-world truth. This dimension is difficult to verify automatically without an external reference source. For this reason, in this project accuracy is interpreted mainly as plausibility.

A plausibility check does not always prove that a value is wrong. Instead, it identifies suspicious values that require further inspection.

Timeliness Timeliness measures whether data are temporally suitable for their intended use. Since this project analyzes historical Formula 1 seasons, timeliness does not mean that the data must be updated to the current day. Instead, it means that dates must be coherent with the temporal scope of the analysis.

Referential Integrity Referential Integrity verifies that every foreign key value in a child table points to an existing row in the parent table.

Assessment Methodology

The Data Quality Assessment is implemented through a Python script connected to the PostgreSQL reconciled database. The script reads the tables, applies deterministic checks, and exports both aggregated scores and detailed issue files.

DA AGGIUNGERE IN FUTURO =_i COSA è OUTPUT E COSA SIGNIFICA

Score Interpretation

The scorecard uses a simple interpretation scale:

Score Range	Status	Meaning
$score \geq 0.95$	Green	The quality level is acceptable. Only minor monitoring is required.
$0.80 \leq score < 0.95$	Yellow	The quality level is acceptable for exploration, but cleaning should be planned.
$score < 0.80$	Red	The quality problem is critical and may block downstream usage.

Main Output Files

The assessment script is expected to produce the following outputs:

Output file	Purpose
<code>dq_scorecard.all_tables.csv</code>	Detailed scorecard with one row for each table, dimension, and check.
<code>dq_scorecard.by_table.csv</code>	Summary score by table, useful to identify the weakest tables.
<code>dq_table_profiles.csv</code>	Basic profiling information such as number of rows, number of columns, and missing values.
<code>dq_issues.lap.csv</code>	Row-level issues detected in the <code>lap</code> table.

Output file	Purpose
<code>dq-issues_result.csv</code>	Row-level issues detected in the <code>result</code> table.
<code>dq-issues_referential_integrity.csv</code>	Broken references between child and parent tables.
<code>dq-scorecard_for_llm.json</code>	Compact JSON version of the scorecard, useful only for generating a natural-language explanation.

Table-by-Table Data Quality Plan

This section defines the planned checks for each table of the reconciled database. The actual numerical results will be inserted after running the Python assessment script.

Season Table

Although the `season` table was manually created from domain knowledge and contains only two records, it is still included in the Data Quality Assessment. The goal is to verify that the manually introduced domain table is complete.

Dimension	Checks
Completeness	<code>season_year</code> must be present because it is the structural identifier of the season
Uniqueness	<code>season_year</code> must be unique, since it is the primary key of the table.
Validity	<code>season_year</code> must belong to the selected analytical scope, namely 2021 or 2022. If present, <code>season_start_date</code> and <code>season_end_date</code> must be valid dates, and <code>number_of_events</code> must be greater than zero.
Consistency	If both dates are present, <code>season_start_date</code> must be before <code>season_end_date</code> . Moreover, the year component of the season dates must be coherent with <code>season_year</code> .
Accuracy	Not implemented since the <code>season</code> table is a manually introduced domain table, its accuracy is documented through source provenance during the re-engineering phase.
Timeliness	Not applicable as a freshness check, because the table represents closed historical seasons. Temporal coherence is already evaluated through validity and consistency checks on <code>season_year</code> and date attributes.

Circuit Table

It is partly based on domain knowledge and includes derived categorical attributes such as circuit and sector categories.

Dimension	Checks
Completeness	<code>circuit_id</code> , <code>global_circuit_category</code> , <code>sector1_category</code> , <code>sector2_category</code> , and <code>sector3_category</code> must be present. These attributes are required because they identify the circuit and support the analytical classification used in the project.

Dimension	Checks
Uniqueness	<code>circuit_id</code> must be unique, since it is the primary key of the table
Validity	<code>global_circuit_category</code> must belong to the allowed global circuit domain. <code>sector1_category</code> , <code>sector2_category</code> , and <code>sector3_category</code> must belong to the allowed sector category domain.
Consistency	Textual attributes such as <code>circuit_name</code> , <code>country</code> , and <code>location</code> are checked for empty strings, duplicated variants, and inconsistent formatting. A geographical consistency check between <code>country</code> and <code>location</code> is not implemented automatically because it would require an external geographical reference source.
Accuracy/Plausibility	Not implemented as a fully automatic check. Their plausibility is documented through the manual classification process.
Timeliness	Not applicable as a freshness check. Possible future circuit layout changes are outside the scope of this assessment.

Allowed domains:

- Global circuit category: `Street`, `PowerSensitive`, `AeroSensitive`, `Mixed`
- Sector category: `Power`, `FastCorners`, `SlowCorners`, `Technical`

Driver Table

The `driver` table contains stable driver master data obtained from the Ergast/Jolpica source through the FastF1 interface.

Dimension	Checks
Completeness	<code>driver_id</code> , <code>abbreviation</code> , and <code>last_name</code> must be present. These attributes are required to identify and describe each driver in the reconciled database and in the later analytical model.
Uniqueness	<code>driver_id</code> must be unique, since it is the primary key of the table.
Validity	If present, <code>date_of_birth</code> must be a valid date and <code>driver_url</code> must have a valid URL format. If present, <code>permanent_number</code> must be a positive numeric value. <code>abbreviation</code> must follow the expected driver-code format, namely a short uppercase alphabetic code.
Consistency	Textual attributes such as <code>first_name</code> , <code>last_name</code> , <code>nationality</code> , and <code>abbreviation</code> are checked for empty strings and inconsistent formatting.
Accuracy/Plausibility	Not implemented as a fully automatic truth check. Plausibility can be assessed by flagging abnormal values, such as implausible driver ages during the selected 2021–2022 seasons.
Timeliness	Not applicable as a freshness check.

Team Table

The **team** table contains constructor/team master data obtained from the Ergast/Jolpica source through the FastF1 interface. It is used as a reference table to identify teams in results and lap-level data.

Dimension	Checks
Completeness	team_id and team_name must be present. team_id is required because it is the structural identifier used in joins, while team_name is required as the readable business label used in reports and visual analyses.
Uniqueness	team_id must be unique, since it is the primary key of the table
Validity	team_id must follow the expected identifier format, namely a non-empty lowercase textual code. If present, team_url must have a valid URL format. Textual attributes such as team_name and nationality must not be empty strings after trimming.
Consistency	In the final reconciled table, textual attributes are checked for duplicated variants, extra spaces, and inconsistent formatting.
Accuracy/Plausibility	Team names and nationalities can be manually verified through source provenance if needed.
Timeliness	Not applicable as a freshness check.

Grand Prix Table

The **grand_prix** table represents a Formula 1 Grand Prix event in a specific season. Each Grand Prix belongs to one season and is associated with one circuit.

Dimension	Checks
Completeness	grand_prix_id , season_year , round_number , and circuit_id must be present. These attributes are required to identify the Grand Prix, connect it to the corresponding season, and associate it with the circuit where it takes place.
Uniqueness	grand_prix_id must be unique, since it is the surrogate primary key of the table. The natural key season_year + round_number should also be unique, because a round number identifies one Grand Prix within a specific season.
Validity	round_number must be greater than zero. If present, event_date must be a valid date. If present, event_format must belong to the allowed event format domain, such as conventional or sprint . Testing events must not be included in the reconciled Grand Prix table. If present, f1_api_support must be a boolean value.
Consistency	If event_date is present, it must be coherent with season_year . Textual attributes such as event_name , official_event_name , country , and location are checked for empty strings, duplicated variants, and inconsistent formatting.

Dimension	Checks
Accuracy/Plausibility	Not implemented as a fully automatic truth check for all event attributes. For this table, the main plausibility control concerns the matching between each Grand Prix and its circuit. Since <code>circuit_id</code> is introduced during the re-engineering phase by matching <code>country</code> and <code>location</code> with the <code>circuit</code> table, the correctness of this association must be validated. Rows with missing or ambiguous matches are flagged for manual review.
Timeliness	Temporal checks are limited to the validity and consistency of <code>event_date</code> with the selected 2021–2022 project scope.
Referential Integrity	<code>season_year</code> must exist in the <code>season</code> table and <code>circuit_id</code> must exist in the <code>circuit</code> table.

Session Table

The `session` table represents a Qualifying or Race session belonging to a specific Grand Prix. Free practice sessions are outside the analytical scope of the project.

Dimension	Checks
Completeness	<code>session_id</code> , <code>grand_prix_id</code> , <code>session_type</code> must be present.
Uniqueness	<code>session_id</code> must be unique. The natural key <code>grand_prix_id</code> + <code>session_type</code> should also be unique.
Validity	<code>session_type</code> must be either Q or R.
Consistency	<code>session_date</code> should be coherent with the corresponding Grand Prix date.
Referential Integrity	<code>grand_prix_id</code> must exist in the <code>grand_prix</code> table.

Result Table

This table is structurally important because it is used to connect drivers and teams to sessions and also supports the enrichment of lap data.

Dimension	Checks
Completeness	<code>result_id</code> , <code>session_id</code> , <code>driver_id</code> , <code>team_id</code> , and <code>position</code> should be present.
Uniqueness	<code>result_id</code> must be unique. The natural key <code>session_id</code> + <code>driver_id</code> should also be unique.
Validity	<code>position</code> must be greater than zero when present. <code>points</code> must be greater than or equal to zero.
Consistency	<code>classified_position</code> , <code>status</code> , <code>time_ms</code> , and <code>laps</code> should be checked for coherent interpretation.
Referential Integrity	<code>session_id</code> , <code>driver_id</code> , and <code>team_id</code> must reference existing rows in <code>session</code> , <code>driver</code> , and <code>team</code> .

Lap Table

Each row represents one lap performed by one driver in one specific session.

Dimension	Checks
Completeness	<code>lap_id</code> , <code>session_id</code> , <code>driver_id</code> , <code>team_id</code> , and <code>lap_number</code> must be present because they identify the lap and connect it to the driver, team, and session context. For analytical completeness, the main performance attributes <code>lap_time_ms</code> , <code>sector1_time_ms</code> , <code>sector2_time_ms</code> , <code>sector3_time_ms</code> , and <code>compound</code> are also checked, because they are required to compare lap and sector performance.
Uniqueness	<code>lap_id</code> must be unique, since it is the surrogate primary key of the table. The natural key <code>session_id + driver_id + lap_number</code> should also be unique, because the same driver cannot have two different laps with the same lap number inside the same session.
Validity	<code>lap_number</code> must be greater than zero. If present, <code>lap_time_ms</code> and sector times must be positive values. Speed attributes such as <code>speed_i1</code> , <code>speed_i2</code> , <code>speed_fl</code> , and <code>speed_st</code> must be positive when present. <code>stint</code> and <code>tyre_life</code> must be non-negative or positive according to their domain. <code>compound</code> must belong to the expected tyre compound domain, and boolean attributes such as <code>fresh_tyre</code> , <code>is_personal_best</code> , <code>deleted</code> , <code>fastfl_generated</code> , and <code>is_accurate</code> must contain valid boolean values.
Consistency	For normal timed laps, the sum of <code>sector1_time_ms</code> , <code>sector2_time_ms</code> , and <code>sector3_time_ms</code> should be coherent with <code>lap_time_ms</code> , allowing a small tolerance for rounding or source precision. Session-time attributes should be coherent with the lap timeline. Pit-related attributes must be interpreted consistently: the presence of <code>pit_out_time_ms</code> or <code>pit_in_time_ms</code> indicates a special lap and may explain missing or abnormal lap and sector times. Deleted or inaccurate laps are not automatically removed, but they are flagged for the later cleaning phase.
Accuracy/Plausibility	Plausibility is assessed by flagging suspicious values, such as extremely low or high lap times, sector times, and speed values. These values are not automatically considered wrong, because Formula 1 laps can be affected by pit stops, Safety Car periods, rain, red flags, deleted laps, or other special racing situations.
Timeliness	Temporal quality is evaluated through the coherence of relative time attributes such as <code>time_ms</code> , <code>lap_start_time_ms</code> , and sector session times within each session. If <code>lap_start_date</code> is present, it should be coherent with the corresponding session or Grand Prix date.
Referential Integrity	<code>session_id</code> , <code>driver_id</code> , and <code>team_id</code> must reference existing rows in the <code>session</code> , <code>driver</code> , and <code>team</code> tables.

Weather Table

Each row represents one weather observation associated with a session and identified by a relative session time.

Dimension	Checks
Completeness	<code>weather_id</code> , <code>session_id</code> , <code>time_ms</code> , <code>air_temp</code> , <code>track_temp</code> , <code>rainfall</code> , and <code>wind_speed</code> must be present. The first three attributes are required to identify each weather observation, connect it to the corresponding session, and place it on the session timeline.
Uniqueness	<code>weather_id</code> must be unique, since it is the surrogate primary key of the table. The natural key <code>session_id + time_ms</code> should also be unique, because a session should not contain two different weather observations with the same timestamp.
Validity	Numerical weather attributes must belong to plausible physical ranges. If present, <code>air_temp</code> and <code>track_temp</code> must be plausible temperature values, <code>humidity</code> must be between 0 and 100, <code>pressure</code> must be positive and within a realistic atmospheric range, <code>rainfall</code> must be boolean, <code>wind_direction</code> must be a valid degree value, and <code>wind_speed</code> must be greater than or equal to zero.
Consistency	<code>time_ms</code> represents elapsed time from the beginning of the session, therefore it must be non-negative and should progress coherently within each <code>session_id</code> . Weather variables are also checked for internal plausibility: for example, <code>track_temp</code> should be reasonably coherent with <code>air_temp</code> , and combinations such as very low humidity with <code>rainfall = true</code> should be flagged as suspicious rather than automatically treated as errors.
Accuracy/Plausibility	Not implemented as a truth check against an external meteorological source. Extreme but possible weather values are flagged for review.
Timeliness	Not applicable as a freshness check.
Referential Integrity	<code>session_id</code> must reference an existing row in the <code>session</code> table.

Track Status Table

The `track_status` table contains time-varying track condition messages recorded during a specific Formula 1 session. Each row represents a change or state of the track at a given elapsed time within the session.

Like the `weather` table, `track_status` is not a direct to-one branch of the Lap Performance attribute tree, because a session can have many track status records over time. However, it is analytically useful because it can be used during ETL to derive lap-level or session-level context attributes.

Dimension	Checks
Completeness	<code>track_status_id</code> , <code>session_id</code> , <code>time_ms</code> , <code>status</code> , and <code>message</code> must be present. The first three attributes are required to identify the record, connect it to a session, and place it on the session timeline. The attributes <code>status</code> and <code>message</code> are also required because they contain the actual track condition information. Without them, the row would not provide useful analytical context.

Dimension	Checks
Uniqueness	<code>track_status_id</code> must be unique, since it is the surrogate primary key of the table. The natural key <code>session_id + time_ms</code> should also be checked, because a session should not contain multiple track status records with the same timestamp unless this is explicitly justified by the source.
Validity	<code>time_ms</code> must be greater than or equal to zero. <code>status</code> must belong to the documented set of track status codes, and <code>message</code> must belong to the documented set of track status messages. Both attributes must not be empty strings after trimming.
Consistency	<code>time_ms</code> represents elapsed time from the beginning of the session, therefore it must progress coherently within each <code>session_id</code> . In addition, <code>status</code> and <code>message</code> must be coherent with each other: the same numerical status code should always correspond to the same textual message. For example, a status code representing clear track conditions should not be associated with a Safety Car or Red Flag message.
Accuracy/Plausibility	Not implemented as a truth check against an external race-control source. Suspicious cases, such as unknown status codes, unknown messages, or inconsistent status-message pairs, are flagged for review.
Timeliness	Not applicable as a freshness check.
Referential Integrity	<code>session_id</code> must reference an existing row in the <code>session</code> table.

Expected status-message domain:

Status code	Message	Meaning
1	AllClear	Normal track condition.
2	Yellow	Yellow flag condition.
4	SCDeployed	Safety Car deployed.
5	Red	Red flag condition.
6	VSCDeployed	Virtual Safety Car deployed.
7	VSCEnding	Virtual Safety Car ending phase.

Planned Visual Outputs

The Python assessment script should also generate visual outputs to summarize the results. These figures will be inserted in the final version of the report after running the script.

Overall Data Quality Score by Table

This plot compares the average data quality score of each table. It is useful to identify which table requires the most cleaning effort.




Figure 1: Overall Data Quality Score by Table

Data Quality Score by Dimension

This plot compares the quality dimensions across the whole database. It is useful to understand whether the main problem is completeness, uniqueness, validity, consistency, timeliness, accuracy, or referential integrity.




Figure 2: Data Quality Score by Dimension

Heatmap of Table-Dimension Scores

A heatmap is useful to show which tables and dimensions have the weakest scores. This figure allows the reader to identify problematic combinations, such as low completeness in **lap** or low validity in **weather**.




Figure 3: Heatmap of Data Quality Scores by Table and Dimension

Missing Values Summary

This plot summarizes missing values by table or by column. It supports the completeness analysis and helps identify attributes that require special attention during cleaning.




Figure 4: Missing Values Summary

Referential Integrity Issues

This plot summarizes broken references between child and parent tables. It is especially important before applying explicit SQL foreign key constraints.




Figure 5: Referential Integrity Issues by Relationship

Focused Missing Values Analysis

Missing values are treated as a focused data quality problem mainly related to the *Completeness* dimension. However, in this project, not every null value is automatically considered a data quality error. Some attributes may be null because the information is not applicable in a specific Formula 1 context.

For this reason, the missing values analysis does not simply count null values. It also classifies them according to their meaning in the domain.

Missing Values and Completeness

The first interpretation of missing values is connected to completeness. A table is complete when all required values are present. In this project, structural identifiers and foreign keys are considered mandatory, because they are necessary to connect the reconciled tables.

Examples of structural attributes are:

- `session_id`
- `driver_id`
- `team_id`
- `lap_id`
- `result_id`

If one of these attributes is missing, the row may become unusable for relational joins and for the later Data Warehouse loading phase.

By contrast, some analytical attributes may be missing without immediately invalidating the row. For example, a missing `lap_time_ms` must be investigated, but it is not equivalent to a missing primary or foreign key.

MCAR, MAR, and MNAR Classification

The theoretical classification of missing values distinguishes between three main mechanisms: MCAR, MAR, and MNAR.

Mechanism	Meaning	Interpretation in this project
MCAR	Missing Completely At Random. The missingness has no relationship with other variables in the dataset.	Possible case for small random gaps in contextual data, such as occasional missing weather observations.
MAR	Missing At Random. The missingness depends on other observed variables.	Possible case for qualifying times, where <code>q2_ms</code> or <code>q3_ms</code> may be missing depending on the driver's progression through the qualifying session.
MNAR	Missing Not At Random. The missingness depends on the missing value itself or on an unobserved reason.	Possible suspicious case when missing lap times are related to problematic laps, deleted laps, inaccurate laps, or special racing conditions.

This classification is useful because different missing mechanisms require different cleaning strategies. In this project, missing values are not automatically imputed. Instead, they are first classified and then used to define motivated cleaning or filtering decisions.

Expected NULL Values

Some null values are expected and should not be interpreted as data quality errors. These values are null because the corresponding attribute is meaningful only under specific conditions.

Examples include:

Attribute	Condition	Interpretation
<code>q1_ms</code> , <code>q2_ms</code> , <code>q3_ms</code>	Session type is Race	Expected null, because qualifying segment times are not applicable to race sessions.
<code>q2_ms</code>	Driver did not reach Q2	Expected null, because the driver has no Q2 time.
<code>q3_ms</code>	Driver did not reach Q3	Expected null, because the driver has no Q3 time.
<code>pit_in_time_ms</code>	Normal lap	Expected null, because the lap is not an in-lap.
<code>pit_out_time_ms</code>	Normal lap	Expected null, because the lap is not an out-lap.
<code>deleted_reason</code>	Lap is not deleted	Expected null, because no deletion reason exists.

Therefore, the missing values analysis must distinguish between expected nulls and real data quality problems.

Problematic NULL Values

A null value is considered problematic when it affects the identification, connection, or analytical usability of a row.

Examples of problematic null values are:

- missing primary keys or surrogate identifiers;
- missing foreign keys required to join tables;
- missing `driver_id` or `team_id` in fact-like tables;
- missing `session_id` in session-dependent tables;
- missing `lap_number` in the lap table;
- missing `lap_time_ms` in laps that should represent valid timed laps.

These missing values are flagged by the assessment script and are later considered during the cleaning phase.

The output of this analysis is expected to include:

- a summary of missing values by table;
- a summary of missing values by column;
- a classification of expected and problematic null values;
- a list of rows requiring further inspection.

Focused Outlier Detection

Outlier detection is treated as a focused data quality analysis related mainly to *Accuracy*, *Validity*, and *Consistency*. The goal is to identify values that are statistically unusual or suspicious from a domain perspective.

In this project, an outlier is not automatically considered a wrong value. Formula 1 data can contain extreme but valid values due to pit stops, Safety Car periods, red flags, weather conditions, deleted laps, inaccurate laps, or other special racing situations.

Therefore, outlier detection is used to flag suspicious records, not to automatically delete them.

IQR Method

The Interquartile Range method is a simple and interpretable univariate outlier detection technique. It is based on the first quartile, the third quartile, and the interquartile range.

$$IQR = Q3 - Q1$$

A value is considered suspicious if it falls outside the following interval:

$$LowerBound = Q1 - k \cdot IQR$$

$$UpperBound = Q3 + k \cdot IQR$$

where k is usually set to 1.5 for standard outlier detection or to a larger value for a more conservative analysis.

In this project, the IQR method can be applied to numerical attributes such as:

- `lap_time_ms`

- `sector1_time_ms`
- `sector2_time_ms`
- `sector3_time_ms`
- speed attributes
- weather attributes

Modified Z-score

The Modified Z-score is a robust alternative to the classical Z-score. It uses the median and the Median Absolute Deviation instead of the mean and standard deviation.

$$M_i = \frac{0.6745(x_i - \text{median}(x))}{MAD}$$

where:

$$MAD = \text{median}(|x_i - \text{median}(x)|)$$

A value is usually considered suspicious when:

$$|M_i| > 3.5$$

This method is useful because Formula 1 performance data may be skewed. For example, lap times may include normal fast laps, slow laps, pit laps, laps under Safety Car, and other special cases. Using the median makes the method less sensitive to extreme values.

Consensus and Multi-method Flag

No single outlier detection method is perfect. For this reason, the assessment uses a simple consensus strategy.

Each numerical value can be flagged by one or more methods. The number of methods that flag the same value is used as an indicator of anomaly strength.

Number of Flags	Interpretation	Action
0	Normal value	No action required.
1	Weak anomaly	Review only if the attribute is important for the analysis.
2 or more	Stronger suspicious value	Investigate before using the row in the Data Warehouse.

In the first implementation, the consensus is based on the IQR method and the Modified Z-score. More advanced techniques, such as Isolation Forest or Local Outlier Factor, can be added later if necessary.

Domain Interpretation

The statistical detection of an outlier is not sufficient to decide whether a record must be removed or corrected. Each suspicious value must be interpreted in the Formula 1 domain.

For example, a very slow lap can be caused by:

- an in-lap;

- an out-lap;
- a pit stop;
- a Safety Car period;
- a Virtual Safety Car period;
- rain or wet track conditions;
- red flags or yellow flags;
- an inaccurate or deleted lap.

For this reason, the output of the outlier detection script is used as input for later cleaning and analytical filtering decisions. The script does not delete outliers automatically. It only flags them and provides evidence for manual or rule-based interpretation.

The expected outputs are:

- an outlier summary by table and column;
- row-level outlier flags;
- a consensus score for suspicious values;
- plots showing outlier distributions;
- a list of records requiring further review.

Optional LLM Support

A Large Language Model can be used as an optional support layer for interpreting the Data Quality Assessment outputs. However, it is not part of the deterministic data quality pipeline and it does not replace the Python-based checks.

The official assessment results are produced only by deterministic scripts. These scripts compute quality scores, detect missing values, identify outliers, and generate row-level issue files. The LLM can only be used after these outputs have been generated.

Not Part of the Deterministic Pipeline

The LLM is not used to decide whether a value is valid, whether a row must be removed, or whether a cleaning rule must be applied.

The deterministic part of the pipeline includes:

- general data quality checks;
- missing values analysis;
- outlier detection;
- referential integrity checks;
- row-level flags;
- scorecard generation.

These steps are reproducible because the same input data and the same rules always produce the same outputs.

By contrast, LLM outputs are probabilistic and may vary depending on the prompt, model, and configuration. For this reason, the LLM is separated from the official assessment logic.

Used Only for Explanation

If used, the LLM receives only aggregated and already generated outputs, such as:

- data quality scorecards;
- missing value summaries;
- missing value classifications;
- outlier summaries;
- referential integrity issue summaries.

The goal is to generate a natural-language explanation of the results. Possible outputs include:

- a stakeholder-oriented data quality summary;
- a technical explanation of the main quality problems;
- a draft paragraph for the final report;
- a preliminary list of cleaning priorities.

The LLM does not access the PostgreSQL database directly and does not modify any data.

Manually Validated Output

Every LLM-generated output must be manually reviewed before being included in the project documentation.

The LLM output is treated as a draft, not as an authoritative result. The final explanation included in the report must be checked against the original scorecards and issue files.

For transparency, the following files can be stored:

- `llm_input_dqa_summary.json`
- `llm_prompt_dqa_interpretation.txt`
- `llm_output_dqa_interpretation.md`
- `llm_human_review_notes.md`

This makes the use of the LLM explicit, controlled, and auditable.

In summary, the LLM is used only as an explanatory assistant. The official data quality decisions remain based on deterministic and reproducible checks.

Conclusion

The final goal is not only to detect problems, but also to produce evidence for the next phase. The generated scorecards, issue files, and plots will guide the cleaning script and will support the documentation of all data quality decisions in the final project report.